

Experiment – 07

Demonstrate the use of Sqoop tools to transfer data between Hadoop & relation database servers

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ mysql -u root -pcloudera  
Welcome to the MySQL monitor.  Commands end with ; or \g.  
Your MySQL connection id is 17  
Server version: 5.1.66 Source distribution  
  
Copyright (c) 2000, 2012, Oracle and/or its affiliates. All rights reserved.  
  
Oracle is a registered trademark of Oracle Corporation and/or its  
affiliates. Other names may be trademarks of their respective  
owners.  
  
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.  
mysql>
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ mysql -u root -pcloudera  
Welcome to the MySQL monitor.  Commands end with ; or \g.  
Your MySQL connection id is 17  
Server version: 5.1.66 Source distribution  
  
Copyright (c) 2000, 2012, Oracle and/or its affiliates. All rights reserved.  
  
Oracle is a registered trademark of Oracle Corporation and/or its  
affiliates. Other names may be trademarks of their respective  
owners.  
  
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.  
mysql> show databases;  
+-----+  
| Database |  
+-----+  
| information_schema |  
| cm |  
| firehose |  
| hue |  
| metastore |  
| mysql |  
| nav |  
| navms |  
| oozie |  
| retail_db |  
| rman |  
| sentry |  
+-----+  
12 rows in set (0.70 sec)  
  
mysql>
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
+-----+  
12 rows in set (0.70 sec)  
  
mysql> use retail db;  
Reading table information for completion of table and column names  
You can turn off this feature to get a quicker startup with -A  
  
Database changed  
mysql> show tables;  
+-----+  
| Tables_in_retail_db |  
+-----+  
| categories           |  
| customers            |  
| departments          |  
| order_items         |  
| orders              |  
| products             |  
+-----+  
6 rows in set (0.00 sec)  
  
mysql> select *from departments;  
+-----+-----+  
| department_id | department_name |  
+-----+-----+  
| 2             | Fitness         |  
| 3             | Footwear        |  
| 4             | Apparel         |  
| 5             | Golf            |  
| 6             | Outdoors        |  
| 7             | Fan Shop        |  
+-----+-----+  
6 rows in set (0.00 sec)  
  
mysql> █
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ hostname -f  
quickstart.cloudera  
[cloudera@quickstart ~]$ sqoop list-databases --connect jdbc:mysql://quickstart:3306/ --password cloudera --username root;  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
21/07/11 11:00:40 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.2  
21/07/11 11:00:40 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
21/07/11 11:00:44 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
information_schema  
cm  
firehose  
hue  
metastore  
mysql  
nav  
navms  
oozie  
retail_db  
rman  
sentry  
[cloudera@quickstart ~]$ █
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ sqoop list-databases --connect jdbc:mysql://quickstart:3306/ --password cloudera --username root;  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
21/07/11 11:00:40 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.2  
21/07/11 11:00:40 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
21/07/11 11:00:44 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
information_schema  
cm  
firehose  
hue  
metastore  
mysql  
nav  
navms  
oozie  
retail_db  
rman  
sentry  
[cloudera@quickstart ~]$ sqoop list-tables --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root;  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
21/07/11 11:06:00 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.2  
21/07/11 11:06:01 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
21/07/11 11:06:04 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
categories  
customers  
departments  
order_items  
orders  
products  
[cloudera@quickstart ~]$
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
orders  
products  
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table departments;  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
21/07/11 11:09:21 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.2  
21/07/11 11:09:21 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
21/07/11 11:09:22 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
21/07/11 11:09:22 INFO tool.CodeGenTool: Beginning code generation  
21/07/11 11:09:24 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `departments` AS t LIMIT 1  
21/07/11 11:09:24 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `departments` AS t LIMIT 1  
21/07/11 11:09:25 INFO orm.CompilationManager: HADOOP MAPRED_HOME is /usr/lib/hadoop-mapreduce  
Note: /tmp/sqoop-cloudera/compile/2d50096d5553cb411a3c8007b20013c/departments.java uses or overrides a deprecated API.  
Note: Recompile with -Xlint:deprecation for details.  
21/07/11 11:09:37 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/2d50096d5553cb411a3c8007b20013c/departments.jar  
21/07/11 11:09:38 WARN manager.MySQLManager: It looks like you are importing from mysql.  
21/07/11 11:09:38 WARN manager.MySQLManager: This transfer can be faster! Use the --direct  
21/07/11 11:09:38 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.  
21/07/11 11:09:38 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)  
21/07/11 11:09:38 INFO mapreduce.ImportJobBase: Beginning import of departments  
21/07/11 11:09:38 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address  
21/07/11 11:09:42 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar  
21/07/11 11:09:48 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps  
21/07/11 11:09:49 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032  
21/07/11 11:10:10 INFO db.DBInputFormat: Using read committed transaction isolation  
21/07/11 11:10:10 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`department_id`), MAX(`department_id`) FROM `departments`  
21/07/11 11:10:10 INFO mapreduce.JobSubmitter: number of splits:4  
21/07/11 11:10:13 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1626024224133_0001  
21/07/11 11:10:16 INFO impl.YarnClientImpl: Submitted application application_1626024224133_0001  
21/07/11 11:10:17 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1626024224133_0001/  
21/07/11 11:10:17 INFO mapreduce.Job: Running job: job_1626024224133_0001
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Launched map tasks=4  
Other local map tasks=4  
Total time spent by all maps in occupied slots (ms)=515975  
Total time spent by all reduces in occupied slots (ms)=0  
Total time spent by all map tasks (ms)=515975  
Total vcore-seconds taken by all map tasks=515975  
Total megabyte-seconds taken by all map tasks=528358400  
Map-Reduce Framework  
Map input records=6  
Map output records=6  
Input split bytes=481  
Spilled Records=0  
Failed Shuffles=0  
Merged Map outputs=0  
GC time elapsed (ms)=3548  
CPU time spent (ms)=7940  
Physical memory (bytes) snapshot=389091328  
Virtual memory (bytes) snapshot=6020833280  
Total committed heap usage (bytes)=119799808  
File Input Format Counters  
Bytes Read=0  
File Output Format Counters  
Bytes Written=60  
21/07/11 11:14:34 INFO mapreduce.ImportJobBase: Transferred 60 bytes in 285.7716 seconds (0.21 bytes/sec)  
21/07/11 11:14:34 INFO mapreduce.ImportJobBase: Retrieved 6 records.  
[cloudera@quickstart ~]$ hadoop fs -ls  
Found 3 items  
drwxr-xr-x - cloudera cloudera 0 2021-07-10 11:42 .Trash  
drwxr-xr-x - cloudera cloudera 0 2021-07-10 11:28 Training  
drwxr-xr-x - cloudera cloudera 0 2021-07-11 11:14 departments  
[cloudera@quickstart ~]$
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
21/07/11 11:14:34 INFO mapreduce.ImportJobBase: Retrieved 6 records.  
[cloudera@quickstart ~]$ hadoop fs -ls  
Found 3 items  
drwxr-xr-x - cloudera cloudera 0 2021-07-10 11:42 .Trash  
drwxr-xr-x - cloudera cloudera 0 2021-07-10 11:28 Training  
drwxr-xr-x - cloudera cloudera 0 2021-07-11 11:14 departments  
[cloudera@quickstart ~]$ hadoop fs -ls departments;  
Found 5 items  
-rw-r--r-- 1 cloudera cloudera 0 2021-07-11 11:14 departments/_SUCCESS  
-rw-r--r-- 1 cloudera cloudera 21 2021-07-11 11:14 departments/part-m-00000  
-rw-r--r-- 1 cloudera cloudera 10 2021-07-11 11:14 departments/part-m-00001  
-rw-r--r-- 1 cloudera cloudera 7 2021-07-11 11:14 departments/part-m-00002  
-rw-r--r-- 1 cloudera cloudera 22 2021-07-11 11:14 departments/part-m-00003  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-0000  
cat: '/user/cloudera/departments/part-m-0000': No such file or directory  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-00000  
2,Fitness  
3,Footwear  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-00001  
4,Apparel  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-00002  
5,Golf  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-00003  
6,Outdoors  
7,Fan Shop  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part*  
2,Fitness  
3,Footwear  
4,Apparel  
5,Golf  
6,Outdoors  
7,Fan Shop  
[cloudera@quickstart ~]$
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table departments --  
target-dir /user/cloudera/department1  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
21/07/11 11:24:28 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.2  
21/07/11 11:24:28 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
21/07/11 11:24:30 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
21/07/11 11:24:30 INFO tool.CodeGenTool: Beginning code generation  
21/07/11 11:24:32 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `departments` AS t LIMIT 1  
21/07/11 11:24:32 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `departments` AS t LIMIT 1  
21/07/11 11:24:32 INFO orm.CompilationManager: HADOOP MAPRED HOME is /usr/lib/hadoop-mapreduce  
Note: /tmp/sqoop-cloudera/compile/96abb73c9b8f2a328c21f813e57126ed/departments.java uses or overrides a deprecated API.  
Note: Recompile with -Xlint:deprecation for details.  
21/07/11 11:24:43 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/96abb73c9b8f2a328c21f813e57126ed/departments.jar  
21/07/11 11:24:43 WARN manager.MySQLManager: It looks like you are importing from mysql.  
21/07/11 11:24:43 WARN manager.MySQLManager: This transfer can be faster! Use the --direct  
21/07/11 11:24:43 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.  
21/07/11 11:24:43 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)  
21/07/11 11:24:43 INFO mapreduce.ImportJobBase: Beginning import of departments  
21/07/11 11:24:43 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address  
21/07/11 11:24:44 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar  
21/07/11 11:24:48 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps  
21/07/11 11:24:49 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032  
21/07/11 11:24:58 INFO db.DBInputFormat: Using read committed transaction isolation  
21/07/11 11:24:58 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`department_id`), MAX(`department_id`) FROM `departments`  
21/07/11 11:24:58 INFO mapreduce.JobSubmitter: number of splits:4  
21/07/11 11:24:59 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1626024224133_0002  
21/07/11 11:25:02 INFO impl.YarnClientImpl: Submitted application application_1626024224133_0002  
21/07/11 11:25:02 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1626024224133_0002/  
21/07/11 11:25:02 INFO mapreduce.Job: Running job: job_1626024224133_0002  
21/07/11 11:25:33 INFO mapreduce.Job: Job job_1626024224133_0002 running in uber mode : false  
21/07/11 11:25:34 INFO mapreduce.Job: map 0% reduce 0%  
21/07/11 11:27:56 INFO mapreduce.Job: map 50% reduce 0%
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Merged Map outputs=0  
GC time elapsed (ms)=5989  
CPU time spent (ms)=9550  
Physical memory (bytes) snapshot=377458688  
Virtual memory (bytes) snapshot=6020595712  
Total committed heap usage (bytes)=119799808  
File Input Format Counters  
Bytes Read=0  
File Output Format Counters  
Bytes Written=60  
21/07/11 11:28:07 INFO mapreduce.ImportJobBase: Transferred 60 bytes in 198.6761 seconds (0.302 bytes/sec)  
21/07/11 11:28:07 INFO mapreduce.ImportJobBase: Retrieved 6 records.  
[cloudera@quickstart ~]$ hadoop fs -ls  
Found 4 items  
drwxr-xr-x - cloudera cloudera 0 2021-07-10 11:42 .Trash  
drwxr-xr-x - cloudera cloudera 0 2021-07-10 11:28 Training  
drwxr-xr-x - cloudera cloudera 0 2021-07-11 11:28 department1  
drwxr-xr-x - cloudera cloudera 0 2021-07-11 11:14 departments  
[cloudera@quickstart ~]$ hadoop fs -ls /user/cloudera/department1  
Found 5 items  
-rw-r--r-- 1 cloudera cloudera 0 2021-07-11 11:28 /user/cloudera/department1/ SUCCESS  
-rw-r--r-- 1 cloudera cloudera 21 2021-07-11 11:27 /user/cloudera/department1/part-m-00000  
-rw-r--r-- 1 cloudera cloudera 10 2021-07-11 11:27 /user/cloudera/department1/part-m-00001  
-rw-r--r-- 1 cloudera cloudera 7 2021-07-11 11:27 /user/cloudera/department1/part-m-00002  
-rw-r--r-- 1 cloudera cloudera 22 2021-07-11 11:27 /user/cloudera/department1/part-m-00003  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/department1/part*  
2,Fitness  
3,Footwear  
4,Apparel  
5,Golf  
6,Outdoors  
7,Fan Shop  
[cloudera@quickstart ~]$
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table departments --where "department_id=4" --target-dir /user/cloudera/department2;  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
21/07/11 11:39:32 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.2  
21/07/11 11:39:32 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
21/07/11 11:39:33 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
21/07/11 11:39:33 INFO tool.CodeGenTool: Beginning code generation  
21/07/11 11:39:35 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `departments` AS t LIMIT 1  
21/07/11 11:39:35 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `departments` AS t LIMIT 1  
21/07/11 11:39:35 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce  
Note: /tmp/sqoop-cloudera/compile/78f69b1070d3a7e50dd3b946d0b5880e/departments.java uses or overrides a deprecated API.  
Note: Recompile with -Xlint:deprecation for details.  
21/07/11 11:39:46 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/78f69b1070d3a7e50dd3b946d0b5880e/departments.jar  
21/07/11 11:39:46 WARN manager.MySQLManager: It looks like you are importing from mysql.  
21/07/11 11:39:46 WARN manager.MySQLManager: This transfer can be faster! Use the --direct  
21/07/11 11:39:46 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.  
21/07/11 11:39:46 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)  
21/07/11 11:39:46 INFO mapreduce.ImportJobBase: Beginning import of departments  
21/07/11 11:39:46 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address  
21/07/11 11:39:47 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar  
21/07/11 11:39:50 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps  
21/07/11 11:39:51 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032  
21/07/11 11:40:02 INFO db.DBInputFormat: Using read committed transaction isolation  
21/07/11 11:40:02 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`department_id`), MAX(`department_id`) FROM `departments` WHERE ( `department_id`=4 )  
21/07/11 11:40:03 INFO mapreduce.JobSubmitter: number of splits:3  
21/07/11 11:40:05 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1626024224133_0003  
21/07/11 11:40:06 INFO impl.YarnClientImpl: Submitted application application_1626024224133_0003  
21/07/11 11:40:06 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1626024224133_0003/  
21/07/11 11:40:06 INFO mapreduce.Job: Running job: job_1626024224133_0003  
21/07/11 11:40:35 INFO mapreduce.Job: Job job_1626024224133_0003 running in uber mode : false
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Total megabyte-seconds taken by all map tasks=211574784  
Map-Reduce Framework  
  Map input records=3  
  Map output records=3  
  Input split bytes=361  
  Spilled Records=0  
  Failed Shuffles=0  
  Merged Map outputs=0  
  GC time elapsed (ms)=3003  
  CPU time spent (ms)=6810  
  Physical memory (bytes) snapshot=309657600  
  Virtual memory (bytes) snapshot=4515348480  
  Total committed heap usage (bytes)=89849856  
File Input Format Counters  
  Bytes Read=0  
File Output Format Counters  
  Bytes Written=29  
21/07/11 11:41:51 INFO mapreduce.ImportJobBase: Transferred 29 bytes in 121.0245 seconds (0.2396 bytes/sec)  
21/07/11 11:41:51 INFO mapreduce.ImportJobBase: Retrieved 3 records.  
[cloudera@quickstart ~]$  
[cloudera@quickstart ~]$ hadoop fs -ls /user/cloudera/department2  
Found 4 items  
-rw-r--r-- 1 cloudera cloudera      0 2021-07-11 11:41 /user/cloudera/department2/_SUCCESS  
-rw-r--r-- 1 cloudera cloudera      7 2021-07-11 11:41 /user/cloudera/department2/part-m-00000  
-rw-r--r-- 1 cloudera cloudera    11 2021-07-11 11:41 /user/cloudera/department2/part-m-00001  
-rw-r--r-- 1 cloudera cloudera    11 2021-07-11 11:41 /user/cloudera/department2/part-m-00002  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/department2  
cat: '/user/cloudera/department2': Is a directory  
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/department2/part*  
5,Golf  
6,Outdoors  
7,Fan Shop  
[cloudera@quickstart ~]$
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
+-----+  
| Tables_in_retail_db |  
+-----+  
| categories          |  
| customers           |  
| departments         |  
| order_items        |  
| orders             |  
| products            |  
+-----+  
6 rows in set (0.00 sec)  
  
mysql> select *from departments;  
+-----+  
| department_id | department_name |  
+-----+  
| 2             | Fitness         |  
| 3             | Footwear        |  
| 4             | Apparel         |  
| 5             | Golf            |  
| 6             | Outdoors        |  
| 7             | Fan Shop        |  
+-----+  
6 rows in set (0.00 sec)  
  
mysql> create table dpt(department_id int not null auto_increment,department_name varchar(50) not null,primary key(department_id));  
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to use  
near 'auto_increment,department_name varchar(50) not null,primary key(department_id))' at line 1  
mysql> create table dpt(department_id int not null auto_increment,department_name varchar(50) not null,primary key(department_id));  
Query OK, 0 rows affected (0.37 sec)  
  
mysql> select *from dpt;  
Empty set (0.56 sec)  
  
mysql>
```

cloudera-quickstart-vm-5.4.2-0-virtualbox [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Applications Places System

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ sqoop export --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table dpt --export-d  
ir /user/cloudera/department2;  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
21/07/11 11:56:22 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.2  
21/07/11 11:56:22 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
21/07/11 11:56:23 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
21/07/11 11:56:23 INFO tool.CodeGenTool: Beginning code generation  
21/07/11 11:56:24 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'dpt' AS t LIMIT 1  
21/07/11 11:56:25 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'dpt' AS t LIMIT 1  
21/07/11 11:56:25 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce  
Note: /tmp/sqoop-cloudera/compile/178875049af4f22ab26c8f0f3128ac0b/dpt.java uses or overrides a deprecated API.  
Note: Recompile with -Xlint:deprecation for details.  
21/07/11 11:56:32 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/178875049af4f22ab26c8f0f3128ac0b/dpt.jar  
21/07/11 11:56:32 INFO mapreduce.ExportJobBase: Beginning export of dpt  
21/07/11 11:56:32 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address  
21/07/11 11:56:34 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar  
21/07/11 11:56:37 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculat  
ive  
21/07/11 11:56:37 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative  
21/07/11 11:56:37 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps  
21/07/11 11:56:38 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032  
21/07/11 11:56:43 INFO input.FileInputFormat: Total input paths to process : 3  
21/07/11 11:56:43 INFO input.FileInputFormat: Total input paths to process : 3  
21/07/11 11:56:43 INFO mapreduce.JobSubmitter: number of splits:3  
21/07/11 11:56:44 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative  
21/07/11 11:56:44 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1626024224133_0004  
21/07/11 11:56:46 INFO impl.YarnClientImpl: Submitted application application_1626024224133_0004  
21/07/11 11:56:46 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1626024224133_0004/  
21/07/11 11:56:46 INFO mapreduce.Job: Running job: job_1626024224133_0004  
21/07/11 11:57:00 INFO mapreduce.Job: Job job_1626024224133_0004 running in uber mode : false  
21/07/11 11:57:00 INFO mapreduce.Job: map 0% reduce 0%  
21/07/11 11:58:05 INFO mapreduce.Job: map 100% reduce 0%  
21/07/11 11:58:07 INFO mapreduce.Job: Job job_1626024224133_0004 completed successfully  
21/07/11 11:58:08 INFO mapreduce.Job: Counters: 30  
File system Counters  
FILE: Number of bytes read=0  
FILE: Number of bytes written=404274  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=683  
HDFS: Number of bytes written=0  
HDFS: Number of read operations=18  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=0  
Job Counters  
Launched map tasks=3  
Data-local map tasks=3
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
6 rows in set (0.00 sec)  
  
mysql> select *from departments;  
+-----+-----+  
| department_id | department_name |  
+-----+-----+  
| 2 | Fitness |  
| 3 | Footwear |  
| 4 | Apparel |  
| 5 | Golf |  
| 6 | Outdoors |  
| 7 | Fan Shop |  
+-----+-----+  
6 rows in set (0.00 sec)  
  
mysql> create table dpt(department_id int not null auto_increment,department_name varchar(50) not null,primary key(department_id));  
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to use  
near 'auto_increment,department_name varchar(50) not null,primary key(department_id))' at line 1  
mysql> create table dpt(department_id int not null auto_increment,department_name varchar(50) not null,primary key(department_id));  
Query OK, 0 rows affected (0.37 sec)  
  
mysql> select *from dpt;  
Empty set (0.56 sec)  
  
mysql> select *from dpt;  
+-----+-----+  
| department_id | department_name |  
+-----+-----+  
| 6 | Outdoors |  
| 7 | Fan Shop |  
| 5 | Golf |  
+-----+-----+  
3 rows in set (0.00 sec)  
  
mysql> 
```